CONCURRENT BINARY SEARCH TREES: DESIGN AND OPTIMIZATIONS

by

Arunmoezhi Ramachandran

APPROVED BY SUPERVISORY COMMITTEE:

_____

Neeraj Mittal, Chair

_____

Balaji Raghavachari

_____

Venkatesan Subbarayan

_____

Rob F Van Der Wijngaart

To my parents and sister

CONCURRENT BINARY SEARCH TREES: DESIGN AND OPTIMIZATIONS

by

ARUNMOEZHI RAMACHANDRAN, BE, MS

DISSERTATION

Presented to the Faculty of

The University of Texas at Dallas

in Partial Fulfillment

of the Requirements

for the Degree of

DOCTOR OF PHILOSOPHY IN

COMPUTER SCIENCE

THE UNIVERSITY OF TEXAS AT DALLAS

May 2016

ACKNOWLEDGMENTS

Ack page

December 2015

PREFACE

This dissertation was produced in accordance with guidelines which permit the inclusion as part of the dissertation the text of an original paper or papers submitted for publication. The dissertation must still conform to all other requirements explained in the "Guide for the Preparation of Master's Theses and Doctoral Dissertations at The University of Texas at Dallas." It must include a comprehensive abstract, a full introduction and literature review, and a final overall conclusion. Additional material (procedural and design data as well as descriptions of equipment) must be provided in sufficient detail to allow a clear and precise judgment to be made of the importance and originality of the research reported.

It is acceptable for this dissertation to include as chapters authentic copies of papers already published, provided these meet type size, margin, and legibility requirements. In such cases, connecting texts which provide logical bridges between different manuscripts are mandatory. Where the student is not the sole author of a manuscript, the student is required to make an explicit statement in the introductory material to that manuscript describing the student's contribution to the work and acknowledging the contribution of the other author(s). The signatures of the Supervising Committee which precede all other material in the dissertation attest to the accuracy of this statement.

CONCURRENT BINARY SEARCH TREES: DESIGN AND OPTIMIZATIONS

Publication No. _____

Arunmoezhi Ramachandran, PhD
The University of Texas at Dallas, 2016

Supervising Professor: Neeraj Mittal

abstract goes here

TABLE OF CONTENTS

LIST OF FIGURES

LIST OF TABLES

CHAPTER 1

LOCK BASED CONCURRENT BINARY SEARCH TREES

1.1 The Lock-Based Algorithm

We first provide an overview of our algorithm. We then describe the algorithm in more detail and also give its pseudo-code. For ease of exposition, we describe our algorithm assuming no memory reclamation, which can be performed using the well-known technique of hazard pointers (Michael, 2004).

1.1.1 Overview of the Algorithm

Every operation in our algorithm uses *seek* function as a subroutine. The seek function traverses the tree from the root node until it either finds the target key or reaches a non-binary node whose next edge to be followed points to a null node. We refer to the path traversed by the operation during the seek as the *access-path*, and the last node in the access-path as the *terminal node*. The operation then compares the target key with the stored key (the key present in the terminal node). Depending on the result of the comparison and the type of the operation, the operation either terminates or moves to the execution phase. In certain cases in which a key may have moved upward along the access-path, the seek function may have to restart and traverse the tree again; details about restarting are provided later. We now describe the next steps for each of the type of operation one-by-one.

Search: A search operation starts by invoking seek operation. It returns true if the stored key matches the target key and false otherwise.

Insert: An insert operation starts by invoking seek operation. It returns false if the target key matches the stored key; otherwise, it moves to the execution phase. In the execution phase, it attempts to insert the key into the tree as a child node of the last node in the access-path using a CAS instruction. If the instruction succeeds, then the operation returns true; otherwise, it restarts by invoking the seek function again.

Delete: A delete operation starts by invoking seek function. It returns false if the stored key does not match the target key; otherwise, it moves to the execution phase. In the execution phase, it attempts to remove the key stored in the terminal node of the access-path. There are two cases depending on whether the terminal node is a binary node (has two children) or not (has at most one child). In the first case, the operation is referred to as *complex delete operation*. In the second case, it is referred to as *simple delete operation*. In the case of simple delete, the terminal node is removed by changing the pointer at the parent node of the terminal node. In the case of complex delete, the key to be deleted is replaced with the *next largest* key in the tree, which will be stored in the *leftmost node* of the *right subtree* of the terminal node.

1.1.2 Details of the Algorithm

As in most algorithms, to make it easier to handle special cases, we use sentinel keys and sentinel nodes. The structure of an empty tree with only sentinel keys (denoted by ∞_1 and ∞_2 with $\infty_1 < \infty_2$) and sentinel nodes (denoted by $\mathbb{R}$ and $\mathbb{S}$) is shown in Figure 1.1.

Our algorithm, like the one in (Natarajan and Mittal, 2014), operates at edge level. A delete operation obtains ownership of the edges it needs to work on by locking them. To enable locking of an edge, we steal a bit from the child addresses of a node referred to as *lock-flag*. We also steal another bit from the child addresses of a node to indicate that the node is undergoing deletion and will be removed from the tree. We denote this bit by

Figure 1.1: Sentinel keys and nodes ($\infty_1 < \infty_2$)

mark-flag. Finally, to avoid the ABA problem, as in Howley and Jones (Howley and Jones, 2012), we use *unique* null pointers. To that end, we steal yet another bit from the child address, referred to as *null-flag*, and use it to indicate whether the address points to a null or a non-null address. So, when an address changes from a non-null value to a null value, we only set the null-flag and the contents of the address are not otherwise modified. This ensures that all null pointers are unique.

We next describe the details of the seek operation, which is executed by all operations (search as well as modify) after which we describe the details of the execution phase of insert and delete operations.

The Seek Phase

A seek function keeps track of the node in the access-path at which it took the last "right turn" (*i.e.*, it last followed a right edge). Let this "right turn" node be referred to as *anchor node* when the traversal reaches the terminal node. Note that the terminal node is the node whose key matched the target key or whose next child edge is set to a null address. For an illustration, please see Figure 1.2. In the latter case (stored key does not match the target key), it is possible that the key may have moved up in the tree. To ascertain that the seek function did not miss the key because it may have moved up during the traversal, we use the following set of conditions that are *sufficient* (but not necessary) to guarantee that the seek function did not miss the key. First, the anchor node is still part of the tree. (For an

4

illustration, see Figure 1.3) Second, the key stored in the anchor node has not changed since it first encountered the anchor node during the (current) traversal. To check for the above two conditions, we determine whether the anchor node is undergoing deletion by examining it right child edge. We discuss the two cases one-by-one.


Figure 1.2: Nodes in the access path of seek

Figure 1.3: A scenario in which the last right turn node is no longer part of the tree

(a) *Right child edge not marked:* In this case, the anchor node is still part of the tree. We next check whether the key stored in the anchor node has changed. If the key

has not changed, then the seek function returns the results of the (current) traversal, which consists of three addresses: (i) the address of the terminal node, (ii) the address of its parent, and (iii) the null address stored in the child field of the terminal node that caused the traversal to terminate. The last address is required to ensure that an insert operation works correctly (specifically to ascertain that the child field of the terminal node has not undergone any change since the completion of the traversal). We refer to it as the *injection point* of the insert operation. On the other hand, if the key has changed, then the seek function restarts from the root of the tree. A possible optimization is that the seek function restarts only if the target key is now less than the anchor node's key.

(b) *Right child edge marked:* In this case, we compare the information gathered in the current traversal about the anchor node with that in the previous traversal, if one exists. Specifically, if the anchor node of the previous traversal is same as that of the current traversal and the keys found in the anchor node in the two traversals also match, then the seek function terminates, but returns the results of the previous traversal (instead of that of the current traversal). This is because the anchor node was definitely part of the tree during the previous traversal since it was reachable from the root of the tree at the beginning of the current traversal. Otherwise, the seek function restarts from the root of the tree.

For insert and delete operations, we refer to the terminal node as the *target node.*

The Execution Phase of an Insert Operation

In the execution phase, an insert operation creates a new node containing the target key. It then adds the new node to the tree at the injection point using a CAS instruction. For an illustration, see Figure 1.4. If the CAS instruction succeeds, then (the new node becomes a

part of the tree and) the operation terminates; otherwise, the operation restarts from the seek phase. Note that the insert operations are lock-free.

Figure 1.4: An illustration of an insert operation

The Execution Phase of a Delete Operation

The execution of a delete operation starts by checking if the target node is a binary node or not. If it is a binary node, then the delete operation is classified as complex; otherwise it is classified as simple.

For a tree node X, let $X.parent$ denote its parent node, and $X.left$ and $X.right$ denote its left and right child node, respectively. Also, hereafter in this section, let T denote the target node of the delete operation under consideration.

(a) *Simple Delete:* In this case, either $T.left$ or $T.right$ is pointing to a null node. Note that both $T.left$ and $T.right$ may be pointing to null nodes in which case T will be a leaf node. Without loss of generality, assume that $T.right$ is a null node. The removal of T involves locking the following three edges: $\langle T.parent, T \rangle$, $\langle T, T.left \rangle$ and $\langle T, T.right \rangle$. For an illustration, see Figure 1.5.

A lock on an edge is obtained by setting the lock-flag in the appropriate child field of the parent node using a CAS instruction. For example, to lock the edge $\langle X, Y \rangle$, where Y is the left child of X, the lock-flag in the left child of X is set to one. If all the edges are locked successfully, then the operation validates that the key stored in the
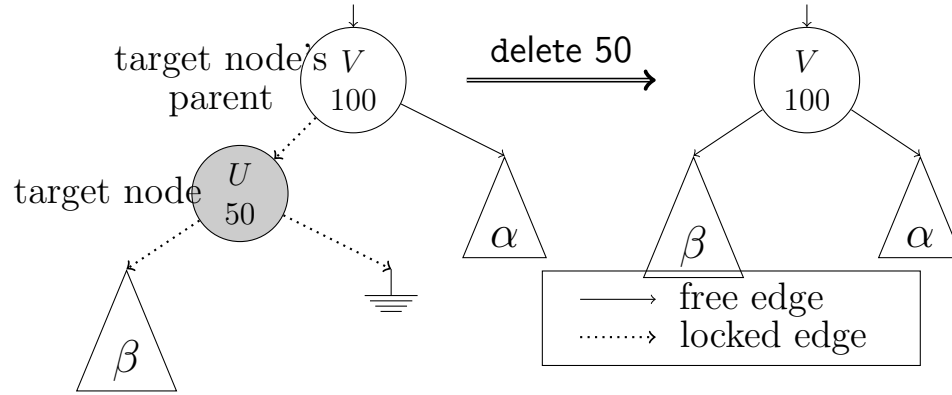
Figure 1.5: An illustration of a simple delete operation

target node still matches the target key. If the validation succeeds, then the operation marks both the children edges of T to indicate that T is going to be removed from the tree. Next, it changes the child pointer at $T.parent$ that is pointing to T to point to $T.left$ using a simple write instruction. Finally, the operation releases all the locks and returns true.

(b) *Complex Delete:* In this case, both $T.left$ and $T.right$ are pointing to non-null nodes. The operation locates the next largest key in the tree, which is the smallest key in the subtree rooted at the right child of T. We refer to this key as the *successor key* and the node storing this key as the *successor node*. Hereafter in this section, let S denote the successor node. Deletion of the key stored in T involves copying the key stored in S to T and then removing S from the tree. To that end, the following edges are locked by setting the lock-flag on the edge using a CAS instruction: $\langle T, T.right\rangle$, $\langle S.parent, S\rangle$, $\langle S, S.left\rangle$ and $\langle S, S.right\rangle$. For an illustration, see Figure 1.6. Note that the first two edges may be same which happens if the successor node is the right child of the target node. Also, since we do not lock the left edge of the target node, the left edge may change and may possibly start pointing to a null address. But, that does not impact the correctness of the complex delete operation.
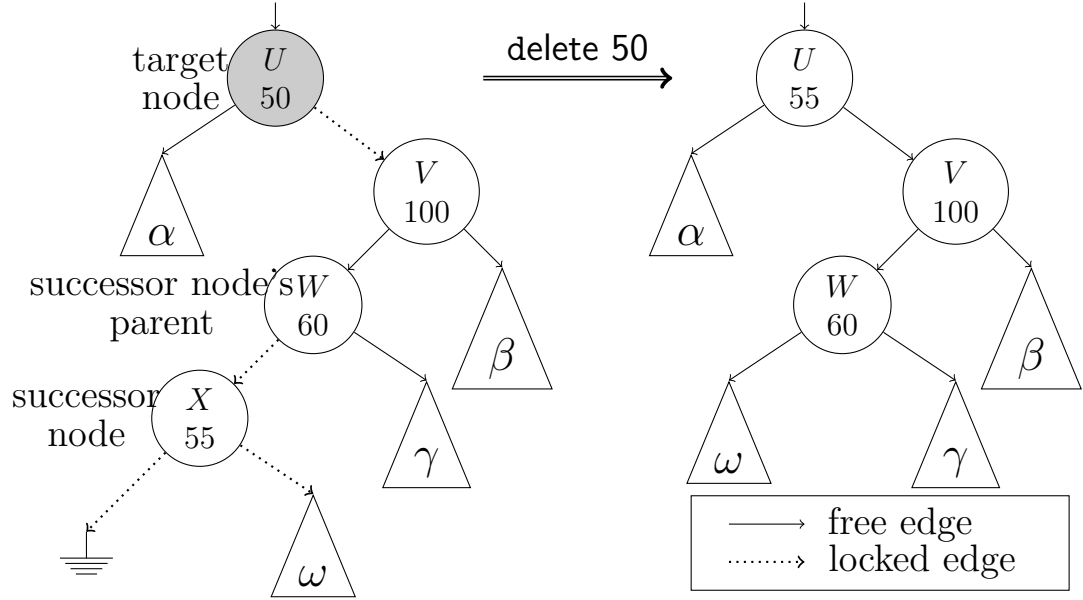
Figure 1.6: An illustration of a complex delete operation

If all the edges are locked successfully, then the operation validates that the key stored in the target node still matches the target key. If the validation succeeds, then the operation copies the key stored in S to T, and marks both the children edges of S to indicate that S is going to be removed from the tree. Next, it changes the child pointer at $S.parent$ that is pointing to S to point to $S.right$ using a simple write instruction. Finally, the operation releases all the locks and returns true.

In both cases (simple as well as complex delete), if the operation fails to obtain any of the locks, then it releases all the locks it was able to acquire up to that point, and restarts from the seek phase. Also, after obtaining all the locks, if the key validation fails, then it implies that some other delete operation has removed the key from the tree while the current execution phase was in progress. In that case, the given delete operation releases all the locks, and simply returns false. Note that using a CAS instruction for setting the lock-flag also enables us to *validate that the child pointer has not changed* since it was last observed in a single step.

Algorithm 1: Data Structures Used

```
   // a tree node
1  struct Node {
2      Key key;
3      { boolean, boolean, boolean, NodePtr } child[2];
       // each child field contains four subfields:  lFlag, mFlag, nFlag and address
4  };

   // used to store the results of a tree traversal
5  struct SeekRecord {
6      NodePtr node;
7      NodePtr parent;
8      NodePtr nullAddress;
9  };

   // used to store information about an anchor node
10 struct AnchorRecord {
11     NodePtr node;
12     Key key;
13 };

   // used to store information about an edge to lock
14 struct LockRecord {
15     NodePtr node;
16     enum { LEFT, RIGHT } which;
17     { boolean, NodePtr } addressSeen;
       // addressSeen contains two subfields:  nFlag and address
18 };

   // local seek record used when looking for a node
19 SeekRecordPtr seekTargetKey, seekSuccessorKey;
   // local array used to store the set of edges to lock
20 LockRecord lockArray[4];
```

1.1.3 Formal Description

We refer to our algorithm as CASTLE (Concurrent Algorithm for Binary Search Tree by Locking Edges).

A pseudo-code of our algorithm is given in Algorithms 1-7. Different data structures used in our algorithm are shown in Algorithm 1. Besides tree node, we use three additional records: (a) *seek record:* to store the outcome of a tree traversal both when looking for the target key and the successor key, (b) *anchor record:* to store information about the anchor node during the seek phase, and (c) *lock record:* to store information about a tree edge that needs to be locked.

Algorithm 2: Seek Function

---

21 SEEK(*key*, *seekRecord*)
22 **begin**
23   **while true do**
        // initialize the variables used in the traversal
24          $pNode := \mathbb{R}; \quad cNode := \mathbb{S};$
25          $address := \mathbb{S} \rightarrow child[\textsf{RIGHT}].address;$
26          $anchorRecord := \{\mathbb{R}, \infty_1\};$
27   **while true do**
            // reached terminal node; read the key stored in the current node
28              $cKey := cNode \rightarrow key;$
29   **if** $key = cKey$ **then**
30                  $seekRecord := \{cNode, pNode, address\};$
31   **return**;

32              $which := key < cKey \ ? \ \textsf{LEFT} : \textsf{RIGHT};$
            // read the next address to dereference along with mark and null flags
33              $\langle *, *, nFlag, address \rangle := cNode \rightarrow child[which];$
34   **if** $nFlag$ **then**   // the null flag is set; reached terminal node
35                  $aNode := anchorRecord \rightarrow node;$
36   **if** $aNode \rightarrow child[\textit{RIGHT}].mFlag$ **then**
                    // the anchor node is marked; it may no longer be part of the tree
37   **if** $anchorRecord = pAnchorRecord$ **then**
                        // the anchor record of the current traversal matches that of
                          the previous traversal
38                          $seekRecord := pSeekRecord;$
39   **return**;
40   **else break**;
41   **else**   // the anchor node is definitely part of the tree
42   **if** $aNode \rightarrow key < key$ **then**   // seek can terminate now
43                      $seekRecord := \{cNode, pNode, address\};$
44   **return**;
45   **else break**;

            // update the anchor record if needed
46   **if** $which = \textit{RIGHT}$ **then**
                // the next edge to be traversed is a right edge; keep track of
                current node and its key
47                  $anchorRecord := \{cNode, cKey\};$

            // traverse the next edge
48              $pNode := cNode; \quad cNode := address;$

Algorithm 3: Search Operation

---

49 **boolean** SEARCH(key)
50 **begin**
51   SEEK(key, $seekTargetKey$);
52      $node := seekTargetKey \rightarrow node$;
53   **if** $node \rightarrow key = key$ **then**
54   **return true**;   `// key found`
55   **else**
56   **return false**;   `// key not found`

Algorithm 4: Insert Operation

---

57 **boolean** INSERT(key)
58 **begin**
59   **while true do**
60   SEEK(key, $seekTargetKey$);
61          $node := seekTargetKey \rightarrow node$;
62   **if** $node \rightarrow key = key$ **then**
63   **return false**;   `// key found`
64   **else**
            `// key not found; add the key to the tree`
65              $newNode :=$ create a new node;
            `// initialize its fields`
66              $newNode \rightarrow key := key$;
67              $newNode \rightarrow child[\textsf{LEFT}] := \langle 0_l, 0_m, 1_n, \textbf{null} \rangle$;
68              $newNode \rightarrow child[\textsf{RIGHT}] := \langle 0_l, 0_m, 1_n, \textbf{null} \rangle$;
            `// determine which child field (left or right) needs to be modified`
69              $which := key < node \rightarrow key$? $\textsf{LEFT}$: $\textsf{RIGHT}$;
            `// fetch the address observed by the seek function in that field`
70              $address := seekTargetKey \rightarrow nullAddress$;
71              $result := \textsf{CAS}(node \rightarrow child[which],$
                                    $\langle 0_l, 0_m, 1_n, address \rangle,$
                                  $\langle 0_l, 0_m, 0_n, newNode \rangle$);
72   **if** $result$ **then**
                `// new key successfully added to the tree`
73   **return true**;

Algorithm 5: Delete Operation

---

74 **boolean** DELETE(*key*)
75 **begin**
76   **while true do**
77   SEEK(*key, seekTargetKey*);
78   *node := seekTargetKey→node*;
79   **if** *node→key ≠ key* **then**   // key not found
80   **return false**;
81   **else**   // key found; read contents of target node's children fields
82   *lField :=* CLEARFLAGS(*node→child*[LEFT]);
83   *rField :=* CLEARFLAGS(*node→child*[RIGHT]);
84   **if** *lField.nFlag*   **or** *rField.nFlag* **then**   // simple delete operation
85   *parent := seekTargetKey→parent*;
86   **if** *key < parent→key* **then** *which :=* LEFT;
87   **else** *which :=* RIGHT;
88   *lockArray*[0] := {*parent, which,* ⟨0, *node*⟩};
89   *lockArray*[1] := {*node,* LEFT, *lField*};
90   *lockArray*[2] := {*node,* RIGHT, *rField*};
91   *result :=* LOCKALL(*lockArray,* 3);
92   **if** *result* **then**   // all locks acquired; perform the operation
93   **if** *node→key = key* **then**   // key still matches; remove the node
94   REMOVECHILD(*parent, which*); *match :=* **true**;
95   **else** *match :=* **false**;
96   UNLOCKALL(*lockArray,* 3);
97   **return** *match*;

98   **else**   // complex delete operation; locate the successor node
99   FINDSMALLEST(*node, rField.address, seekSuccessorKey*);
100   *sNode := seekSuccessorKey→node; sParent := seekSuccessorKey→parent*;
           // determine the edges to be locked
101   *lockArray*[0] := {*node,* RIGHT, *rField*};
102   **if** *node ≠ sParent* **then**
             // successor node is not the right child of target node
103   *lockArray*[1] := {*sParent,* LEFT, ⟨0, *sNode*⟩}; *size :=* 4;
104   **else** *size :=* 3 ;
105   *lField :=* CLEARFLAGS(*sNode→child*[LEFT]);
106   *rField :=* CLEARFLAGS(*sNode→child*[RIGHT]);
107   *lockArray*[*size* − 2] := {*sNode,* LEFT, *lField*};
108   *lockArray*[*size* − 1] := {*sNode,* RIGHT, *rField*};
109   *result :=* LOCKALL(*lockArray, size*);
110   **if** *result* **then**   // all locks acquired; perform the operation
111   **if** *node→key = key* **then**
              // key still matches; copy key in successor node to target node
112   *node→key := sNode→key*;
113   REMOVECHILD(*sParent,* LEFT); *match :=* **true**;
114   **else** *match :=* **false**;
115   UNLOCKALL(*lockArray, size*);
116   **return** *match*;

Algorithm 6: Lock and Unlock Functions

---

117   **boolean** LOCKALL(*lockArray*, *size*)
118   **begin**
119   **for** $i \leftarrow 0$ **to** $size - 1$ **do**
         // acquire lock for the i-th entry
120        $node := lockArray[i].node$;
121        $which := lockArray[i].which$;
122        $lockedAddress := lockArray[i].addressSeen$;
123        $lockedAddress.lFlag := $ **true**;
         // set the lock flag in the child edge
124        $result := \mathsf{CAS}(node \rightarrow child[which], lockArray[i].addressSeen, lockedAddress)$;
125   **if not** *(result)* **then**
           // release all the locks acquired so far
126   UNLOCKALL(*lockArray*, $i - 1$);
127   **return false**;

128   **return true**;

129   UNLOCKALL(*lockArray*, *size*)
130   **begin**
131   **for** $i \leftarrow size - 1$ **to** 0 **do**
132        $node := lockArray[i].node$;
133        $which := lockArray[i].which$;
         // clear the lock flag in the child edge
134        $node \rightarrow child[which].lFlag := $ **false**;

The pseudo-code for the seek function is shown in Algorithm 2. The pseudo-codes for search, insert and delete operations are shown in Algorithm 3, Algorithm 4 and Algorithm 5, respectively. Algorithm 6 contains the pseudo-code for locking and unlocking a set of tree edges, as specified in an array. Finally, Algorithm 7 contains the pseudo-codes for three helper functions used by a delete operation, namely: (a) CLEARFLAGS: to clear lock and mark flags from a child field, (b) FINDSMALLEST: to locate the smallest key in a subtree, and (c) REMOVECHILD: to remove a given child of a node.

In the pseudo-code, to improve clarity, we sometimes use subscripts l, m and n to denote lock, mark and null flags, respectively.

Algorithm 7: Helper Functions used by Delete Operation

---

135   **word** CLEARFLAGS(**word** *field*)
136   **begin**
137   *newField* := *field* with lock and mark flags cleared;
138   **return** *newField*;

139   FINDSMALLEST(*parent, node, seekRecord*)
140   **begin**
     // initialize the variables used in the traversal
141   *pNode* := *parent*;   *cNode* := *node*;
142   **while true do**
143       $\langle *, *, nFlag, address \rangle := cNode \rightarrow child[\text{LEFT}]$;
144   **if not** *(nFlag)* **then**
        // visit the next node
145   *pNode* := *cNode*;   *cNode* := *address*;
146   **else**
        // reached the successor node
147   *seekRecord* := $\{cNode, pNode, address\}$;
148   **break**;

149   REMOVECHILD(*parent, which*)
150   hu **begin**
     // determine the address of the child to be removed
151   *node* := $parent \rightarrow child[which]$;

     // mark both the children edges of the node to be removed
152     $node \rightarrow child[\text{LEFT}].mFlag$:= **true**;
153     $node \rightarrow child[\text{RIGHT}].mFlag$:= **true**;

     // determine whether both the child pointers of the node to be removed are null
154   **if** $node \rightarrow child[\text{LEFT}].nFlag$ **and** $node \rightarrow child[\text{RIGHT}].nFlag$ **then**
      // set the null flag only
155       $parent \rightarrow child[which].nFlag$:= **true**;
156   **else**
      // switch the pointer at the parent to point to its appropriate grandchild
157   **if** $node \rightarrow child[\text{LEFT}].nFlag$ **then**
158   *address* := $node \rightarrow child[\text{RIGHT}].address$;
159   **else** *address* := $node \rightarrow child[\text{LEFT}].address$;
160       $parent \rightarrow child[which].address$:= *address*;

1.1.4 Correctness Proof

It is convenient to treat insert and delete operations that do not change the tree as search operations. We call a tree node *active* if it is reachable from the root of the tree. We call a tree node *passive* if it was active earlier but is not active any more. Note that, before an active node is made passive by a delete operation, both its children edges are *marked*. Also, a CAS instruction performed on an edge (by either an insert operation or a delete operation as part of locking) is successful only if the edge is unmarked. As a result, clearly, if an insert operation completes successfully, then its target node was active when its edge was modified to make the new node (containing the target key) a part of the tree. Likewise, if a delete operation completes successfully, then all the nodes involved in the operation (up to three nodes) were active when their edges were locked.

All Executions are Linearizable

We show that an arbitrary execution of our algorithm is linearizable by specifying the *linearization point* of each operation. Note that the linearization point of an operation is the point during its execution at which the operation appeared to have taken effect. Our algorithm supports three types of operations: search, insert and delete. We now specify the linearization point of each operation.

1. *Insert operation:* The operation is linearized at the point at which it performed the successful CAS instruction that resulted in its target key becoming part of the tree.

2. *Delete operation:* There are two cases depending on whether the delete operation is simple or complex. If the operation is simple delete, then the operation is linearized at the point at which a successful write step was performed at the parent of the target node that resulted in the target node becoming passive. Otherwise, it is linearized at the point at which the original key of the target node was replaced with its successor key.

3. *Search operation:* There are two cases depending on whether the target node was active when the operation read the key stored in the node. If the target node was not active, then the operation is linearized at the point at which the target node became passive. Otherwise, it is linearized at the point at which the read step was performed.

It can be easily verified that, for any execution of the algorithm, the sequence of operations obtained by ordering operations based on their linearization points is legal, *i.e.*, all operations in the sequence satisfy their specification.

Thus we have:

Theorem 1. *Every execution of our algorithm is linearizable.*

All Executions are Deadlock-Free

We say that the system is in a *quiescent state* if no modify operation completes hereafter. We say that the system is in a *potent state* if it has one or more pending modify operations. Note that quiescence is a *stable property*; once the system is in a quiescent state, it stays in a quiescent state. We show that our algorithm is deadlock-free by proving that a potent state is necessarily non-quiescent.

Note that, in a quiescent state, no edges in the tree can be marked. This is because a delete operation marks edges only after it has successfully obtained all the locks, after which it is guaranteed to complete. This also implies that the tree cannot undergo any changes now because that would imply eventual completion of a modify operation. Thus, once a system has reached a quiescent state, all modify operation currently pending repeatedly alternate between seek and execution phases. We say that the system is in a *strongly-quiescent state* if all pending modify operations started their most recent seek phase *after* the system became quiescent. Note that, like quiescence, strong quiescence is also a stable property. Now, once the system has reached a strongly quiescent state, the following can be easily verified. First,

for a given modify operation, every traversal of the tree in the seek phase returns the same target node. Second, for a given delete operation, the set of edges it needs to lock remains the same.

Now, assume that the system eventually reaches a state that is both potent and quiescent. Clearly, from this state, the system will eventually reach a state that is potent and strongly-quiescent. Note that a delete operation in our algorithm locks edges in a *top-down, left-right* manner. As a result, there cannot be a "cycle" involving delete operations. If a delete operation continues to fail in the execution phase, then it is necessarily because it tried to acquire lock on an already locked edge. (Recall that the set of edges does not change any more and there are no marked edges in the tree.) We can construct a chain of operations such that each operation in the chain tried to lock an edge already locked by the next operation in the chain. Clearly, the length of the chain is bounded. This implies that the last operation in the chain is guaranteed to obtain all the locks and will eventually complete. This contradicts the fact that the system is in a quiescent state.

Thus, we have:

Theorem 2. *Every execution of our algorithm is deadlock-free.*

CHAPTER 2

LOCK FREE CONCURRENT BINARY SEARCH TREE

2.1 The Lock-Free Algorithm

For ease of exposition, we describe our algorithm assuming no memory reclamation.

2.1.1 Overview of the Algorithm

Every operation in our algorithm uses *seek* function as a subroutine. The seek function traverses the tree from the root node until it either finds the target key or reaches a non-binary node whose next edge to be followed points to a null node. We refer to the path traversed by the operation during the seek as the *access-path*, and the last node in the access-path as the *terminal node*. The operation then compares the target key with the stored key (the key present in the terminal node). Depending on the result of the comparison and the type of the operation, the operation either terminates or moves to the execution phase. In certain cases in which a key may have moved upward along the access-path, the seek function may have to restart and traverse the tree again; details about restarting are provided later. We now describe the next steps for each of the operations one-by-one.

Search A search operation starts by invoking seek operation. It returns true if the stored key matches the target key and false otherwise.

Insert An insert operation ((shown in Figure 2.2)) starts by invoking seek operation. It returns false if the target key matches the stored key; otherwise, it moves to the execution phase. In the execution phase, it attempts to insert the key into the tree as a child node of

the last node in the access-path using a CAS instruction. If the instruction succeeds, then the operation returns true; otherwise, it restarts by invoking the seek function again.

Delete A delete operation starts by invoking seek function. It returns false if the stored key does not match the target key; otherwise, it moves to the execution phase. In the execution phase, it attempts to remove the key stored in the terminal node of the access-path. There are two cases depending on whether the terminal node is a binary node (has two children) or not (has at most one child). In the first case, the operation is referred to as *complex delete operation*. In the second case, it is referred to as *simple delete operation*. In the case of simple delete (shown in Figure 2.3), the terminal node is removed by changing the pointer at the parent node of the terminal node. In the case of complex delete (shown in Figure 2.4), the key to be deleted is replaced with the *next largest* key in the tree, which will be stored in the *leftmost node* of the *right subtree* of the terminal node.

2.1.2 Details of the Algorithm

As in most algorithms, we use sentinel keys and three sentinel nodes to handle the boundary cases easily. The structure of an empty tree with only sentinel keys (denoted by ∞_0, ∞_1 and ∞_2 with $\infty_0 < \infty_1 < \infty_2$) and sentinel nodes (denoted by $\mathbb{R}$, $\mathbb{S}$ and $\mathbb{T}$) is shown in Figure 2.1.

Our algorithm, like the one in (Natarajan and Mittal, 2014), operates at edge level. A delete operation obtains ownership of the edges it needs to work on by marking them. To enable marking, we steal bits from the child addresses of a node. Specifically, we steal *three* bits from each child address to distinguish between three types of marking: (i) marking for *intent*, (ii) marking for *deletion* and (iii) marking for *promotion*. The three bits are referred to as *intent-flag*, *delete-flag* and *promote-flag*. To avoid the ABA problem, as in Howley and Jones (Howley and Jones, 2012), we use *unique* null pointers. To that end, we steal

Figure 2.1: Nodes in the access path of seek along with sentinel keys and nodes ($\infty_0 < \infty_1 < \infty_2$)

another bit from the child address, referred to as *null-flag*, and use it to indicate whether the address field contains a null or a non-null value. So, when an address changes from a non-null value to a null value, we only set the null-flag and the contents of the address field are not otherwise modified. This ensures that all null pointers are unique.

Finally, we also steal a bit from the key field to indicate whether the key stored in a node is the original key or the replacement key. This information is used in a complex delete operation to coordinate helping among processes.

We next describe the details of the seek function, which is used by all operations (search as well as modify) to traverse the tree after which we describe the details of the execution phase of insert and delete operations.

The Seek Phase

A seek function keeps track of the node in the access-path at which it took the last "right turn" (*i.e.*, it last followed a right edge). Let this "right turn" node be referred to as *anchor node* when the traversal reaches the terminal node. Note that the terminal node is the node whose key matched the target key or whose next child edge is set to a null address. For an illustration, please see Figure 2.1. In the latter case (stored key does not match the target key), it is possible that the key may have moved up in the tree. To ascertain that the seek function did not miss the key because it may have moved up during the traversal, we use the following set of conditions that are *sufficient* (but not necessary) to guarantee that the seek function did not miss the key. First, the anchor node is still part of the tree. Second, the key stored in the anchor node has not changed since it first encountered the anchor node during the (current) traversal. To check for the above two conditions, we determine whether the anchor node is undergoing removal (either delete or promote flag set) by examining its right child edge. We discuss the two cases one-by-one.

(a) *Right child edge not marked:* In this case, the anchor node is still part of the tree. We next check whether the key stored in the anchor node has changed. If the key has not changed, then the seek function returns the results of the (current) traversal, which consists of three addresses: (i) the address of the terminal node, (ii) the address of its parent, and (iii) the null address stored in the child field of the terminal node that caused the traversal to terminate. The last address is required to ensure that an insert operation works correctly (specifically to ascertain that the child field of the terminal node has not undergone any change since the completion of the traversal). We refer

to it as the *injection point* of the insert operation. On the other hand, if the key has changed, then the seek function restarts from the root of the tree.

(b) *Right child edge marked:* In this case, we compare the information gathered in the current traversal about the anchor node with that in the previous traversal, if one exists. Specifically, if the anchor node of the previous traversal is same as that of the current traversal and the keys found in the anchor node in the two traversals also match, then the seek function terminates, but returns the results of the previous traversal (instead of that of the current traversal). This is because the anchor node was definitely part of the tree during the previous traversal since it was reachable from the root of the tree at the beginning of the current traversal. Otherwise, the seek function restarts from the root of the tree.

The seek function also keeps track of the *second-to-last* edge in the access-path (whose endpoints are the parent and grandparent nodes of the terminal node), which is used for helping, if there is a conflict. For insert and delete operations, we refer to the terminal node as the *target node*.

The Execution Phase of an Insert Operation

In the execution phase, an insert operation creates a new node containing the target key. It then adds the new node to the tree at the injection point using a CAS instruction. If the CAS instruction succeeds, then (the new node becomes a part of the tree and) the operation terminates; otherwise, the operation determines if it failed because of a *conflicting* delete operation in progress. If there is no conflicting delete operation in progress, then the operation restarts from the seek phase; otherwise it performs helping and then restarts from the seek phase.

Figure 2.2: An illustration of an insert operation.

The Execution Phase of a Delete Operation

The execution of a delete operation starts in *injection mode*. Once the operation has been injected into the tree, it advances to either *discovery mode* or *cleanup mode* depending on the type of the delete operation.
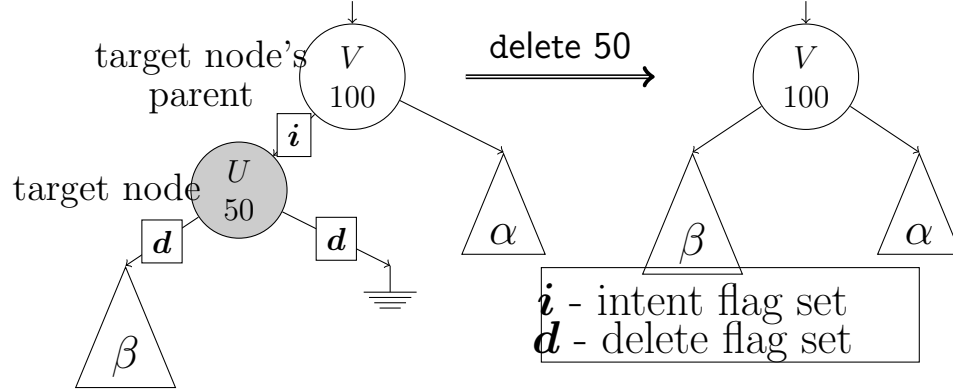
Figure 2.3: An illustration of a simple delete operation.

Injection Mode In the injection mode, the delete operation marks the three edges involving the target node as follows: (i) It first sets the intent-flag on the edge from the parent of the target node to the target node using a CAS instruction. (ii) It then sets the delete-flag on the left edge of the target node using a CAS instruction. (iii) Finally, it sets the delete-flag on the right edge of the target node using a CAS instruction. If the CAS instruction fails at any step, the delete operation performs helping, and either repeats the same step or restarts from the seek phase. Specifically, the delete operation repeats the same step when setting
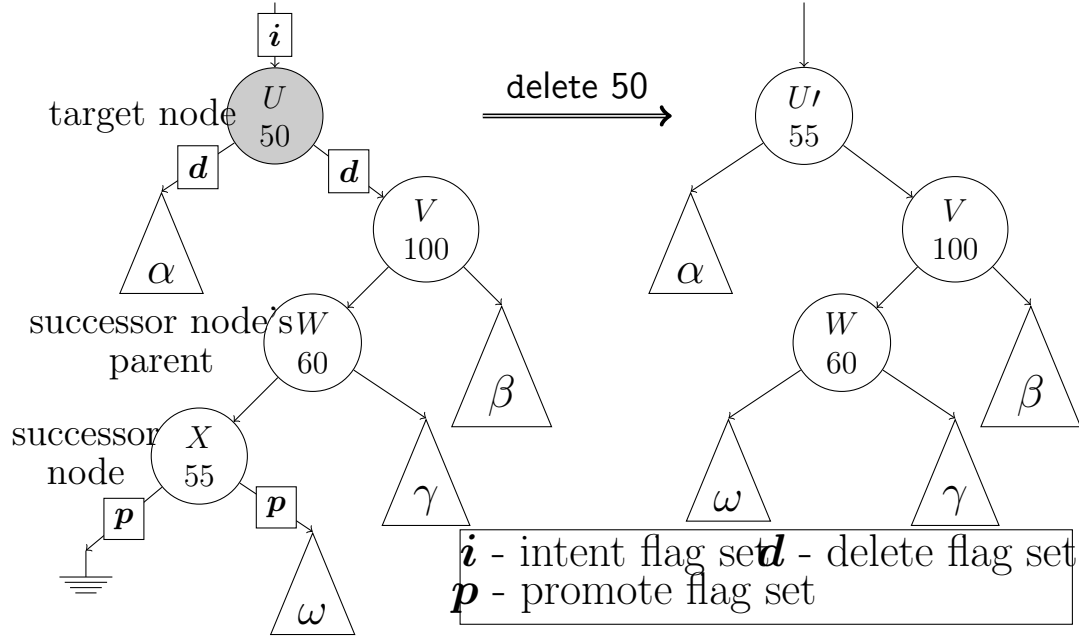
Figure 2.4: An illustration of a complex delete operation.

the delete-flag as long as the target node has not been claimed as the successor node by another delete operation. In all other cases, it restarts from the seek phase.

We maintain the invariant that an edge, once marked, cannot be unmarked. After marking both the edges of the target node, the operation checks whether the target node is a binary node or not. If it is a binary node, then the delete operation is classified as complex; otherwise it is classified as simple. Note that the type of the delete operation cannot change once all the three edges have been marked as described above. If the delete operation is complex, then it advances to the discovery mode after which it will advance to the cleanup mode. On the other hand, if it is simple, then it directly advances to the cleanup mode (and skips the discovery mode). Eventually, the target node is either removed from the tree (if simple delete) or replaced with a "new" node containing the next largest key (if complex delete).

For a tree node X, let $X.parent$ denote its parent node, and $X.left$ and $X.right$ denote its left and right child node, respectively. Also, hereafter in this section, let T denote the target node of the delete operation under consideration.

Discovery Mode   In the discovery mode, a complex delete operation performs the following steps:

1. **Find Successor Key:** The operation locates the next largest key in the tree, which is the smallest key in the subtree rooted at the right child of T. We refer to this key as the *successor key* and the node storing this key as the *successor node*. Hereafter in this section, let S denote the successor node.

2. **Mark Child Edges of Successor Node:** The operation sets the promote-flag on both the child edges of S using a CAS instruction. Note that the left child edge of S will be null. As part of marking the left child edge, we also store the address of T (the target node) in the edge. This is done to enable helping in case the successor node is obstructing the progress of another operation. In case the CAS instruction fails while marking the left child edge, the operation repeats from step 1 after performing helping if needed. On the other hand, if the CAS instruction fails while marking the right child edge, then the marking step is repeated after performing helping if needed.

3. **Promote Successor Key:** The operation replaces the target node's original key with the successor key. At the same time, it also sets the mark bit in the key to indicate that the current key stored in the target node is the replacement key and not the original key.

4. **Remove Successor Node:** The operation removes S (the successor node) by changing the child pointer at $S.parent$ that is pointing to S to point to the right child of S using a CAS instruction. If the CAS instruction succeeds, then the operation advances to the cleanup mode. Otherwise, it performs helping if needed. It then finds S again by performing another traversal of the tree starting from the right child of T. If the traversal

fails to find S (recall that the left edge of S is marked for promotion and contains the address of T), then S has already been removed from the tree by another operation as part of helping, and the delete operation advances to the cleanup mode. On advancing to the cleanup mode, the operation sets a flag in T indicating that S has been removed from the tree (and T can now be replaced with a new node) so that other operations trying to help it know not to look for S.

Algorithm 8: Data Structures Used

```
161  struct Node {
162      {Boolean, Key} mKey;
163      {Boolean, Boolean, Boolean, Boolean, NodePtr} child[2];
164      Boolean readyToReplace;
165  };
166  struct Edge {
167      NodePtr parent, child;
168      enum which { LEFT, RIGHT };
169  };
170  struct SeekRecord {
171      Edge lastEdge, pLastEdge, injectionEdge;
172  };

173  struct AnchorRecord {
174      NodePtr node;
175      Key key;
176  };

177  struct StateRecord {
178      Edge targetEdge, pTargetEdge;
179      Key targetKey, currentKey;
180      enum mode { INJECTION, DISCOVERY, CLEANUP };
181      enum type { SIMPLE, COMPLEX } ;
         // the next field stores pointer to a seek record; it is used for finding the
             successor if the delete operation is complex
182      SeekRecordPtr successorRecord;
183  };
     // object to store information about the tree traversal when looking for a given key
         (used by the seek function)
184  SeekRecordPtr targetRecord := new seek record;
     // object to store information about process' own delete operation
185  StateRecordPtr myState := new state;
```

Algorithm 9: Seek Function

186   SEEK(*key*, *seekRecord*)
187   **begin**
188     $pAnchorRecord := \{\mathbb{S}, \infty_1\};$
189   **while true do**
      // initialize all variables used in traversal
190       $pLastEdge := \{\mathbb{R}, \mathbb{S}, \mathsf{RIGHT}\};$ $lastEdge := \{\mathbb{S}, \mathbb{T}, \mathsf{RIGHT}\};$
191       $curr := \mathbb{T};$ $anchorRecord := \{\mathbb{S}, \infty_1\};$

192   **while true do**
        // read the key stored in the current node
193         $\langle *, cKey \rangle := curr \rightarrow mKey;$
        // find the next edge to follow
194         $which := key < cKey \, ? \, \mathsf{LEFT} \colon \mathsf{RIGHT};$
195         $\langle n, *, d, p, next \rangle := curr \rightarrow child[which];$
        // check for the completion of the traversal
196   **if** $key = cKey$ **or** n **then**
          // either key found or no next edge to follow; stop the traversal
197           $seekRecord \rightarrow pLastEdge := pLastEdge;$
198           $seekRecord \rightarrow lastEdge := lastEdge;$
199           $seekRecord \rightarrow injectionEdge := \{curr, next, which\};$

200   **if** $key = cKey$ **then** // keys match
201   **return;**
202   **else break;**

203   **if** $which = \mathsf{RIGHT}$ **then**
          // next edge to be traversed is a right edge; keep track of the
           current node and its key
204           $anchorRecord := \langle curr, cKey \rangle;$

        // traverse the next edge
205         $pLastEdge := lastEdge;$ $lastEdge := \{curr, next, which\};$ $curr := next;$
      // key was not found; check if can stop
206       $\langle *, *, d, p, * \rangle := anchorRecord.node \rightarrow child[\mathsf{RIGHT}];$
207   **if not** *(d)* **and not** *(p)* **then**
        // anchor node is still part of the tree; check if anchor node's key has
          changed
208         $\langle *, aKey \rangle := anchorRecord.node \rightarrow mKey;$
209   **if** $anchorRecord.key = aKey$ **then return;**
210   **else**
        // check if the anchor record (the node and its key) matches that of the
          previous traversal
211   **if** $pAnchorRecord = anchorRecord$ **then**
          // return the results of the previous traversal
212           $seekRecord := pSeekRecord;$
213   **return;**

      // store the results of the traversal and restart
214       $pSeekRecord := seekRecord;$ $pAnchorRecord := anchorRecord;$

Algorithm 10: Search Operation

---

215 **Boolean** SEARCH(*key*)

216 **begin**

217   SEEK(*key, mySeekRecord*);

218    $node := mySeekRecord \rightarrow lastEdge.child$;

219    $\langle *, nKey \rangle := node \rightarrow mKey$;

220   **if** $nKey = key$ **then return true**;

221   **else return false**;

Algorithm 11: Insert Operation

---

222 **Boolean** INSERT(*key*)

223 **begin**

224   **while true do**

225   SEEK(*key, targetRecord*);

226       $targetEdge := targetRecord \rightarrow lastEdge$;

227       $node := targetEdge.child$;

228       $\langle *, nKey \rangle := node \rightarrow mKey$;

229   **if** $key = nKey$ **then return false**;

          // create a new node and initialize its fields

230       $newNode :=$ create a new node;

231       $newNode \rightarrow mKey := \langle 0_m, key \rangle$;

232       $newNode \rightarrow child[\mathsf{LEFT}] := \langle 1_n, 0_i, 0_d, 0_p, \mathbf{null} \rangle$;

233       $newNode \rightarrow child[\mathsf{RIGHT}] := \langle 1_n, 0_i, 0_d, 0_p, \mathbf{null} \rangle$;

234       $newNode \rightarrow readyToReplace := \mathbf{false}$;

235       $which := targetRecord \rightarrow injectionEdge.which$;

236       $address := targetRecord \rightarrow injectionEdge.child$;

237       $result := \mathsf{CAS}(node \rightarrow child[which], \langle 1_n, 0_i, 0_d, 0_p, address \rangle, \langle 0_n, 0_i, 0_d, 0_p, newNode \rangle)$;

238   **if** *result* **then return true**;

          // help if needed

239       $\langle *, *, d, p, * \rangle := node \rightarrow child[which]$;

240   **if** d **then** HELPTARGETNODE(*targetEdge*) ;

241   **else if** p **then** HELPSUCCESSORNODE(*targetEdge*) ;

Algorithm 12: Delete Operation

242 **Boolean** Delete(*key*)

243 **begin**

// initialize the state record

244 $myState{\rightarrow}targetKey := key;$ $myState{\rightarrow}currentKey := key;$

245 $myState{\rightarrow}mode :=$ INJECTION;

246 **while true do**

247 Seek($myState{\rightarrow}currentKey, targetRecord$);

248 $targetEdge := targetRecord{\rightarrow}lastEdge;$ $pTargetEdge := targetRecord{\rightarrow}pLastEdge;$

249 $\langle *, nKey \rangle := targetEdge.child{\rightarrow}mKey;$

250 **if** $myState{\rightarrow}currentKey \neq nKey$ **then**

// the key does not exist in the tree

251 **if** $myState{\rightarrow}mode =$ *INJECTION* **then return false**;

252 **else return true**;

// perform appropriate action depending on the mode

253 **if** $myState{\rightarrow}mode =$ *INJECTION* **then**

// store a reference to the target edge

254 $myState{\rightarrow}targetEdge := targetEdge;$

255 $myState{\rightarrow}pTargetEdge := pTargetEdge;$

// attempt to inject the operation at the node

256 Inject($myState$);

// mode would have changed if injection was successful

257 **if** $myState{\rightarrow}mode \neq$ *INJECTION* **then**

// check if the target node found by the seek function matches the one stored in the state record

258 **if** $myState{\rightarrow}targetEdge.child \neq targetEdge.child$ **then return true**;

// update the target edge information using the most recent seek

259 $myState{\rightarrow}targetEdge := targetEdge;$

260 **if** $myState{\rightarrow}mode =$ *DISCOVERY* **then**

// complex delete operation; locate the successor node and mark its child edges with promote flag

261 FindAndMarkSuccessor($myState$);

262 **if** $myState{\rightarrow}mode =$ *DISCOVERY* **then**

// complex delete operation; promote the successor node's key and remove the successor node

263 RemoveSuccessor($myState$);

264 **if** $myState{\rightarrow}mode =$ *CLEANUP* **then**

// either remove the target node (simple delete) or replace it with a new node with all fields unmarked (complex delete)

265 $result :=$ Cleanup($myState$);

266 **if** $result$ **then return true**;

267 **else**

268 $\langle *, nKey \rangle := targetEdge.child{\rightarrow}mKey;$

269 $myState{\rightarrow}currentKey := nKey;$

Algorithm 13: Injecting a Deletion Operation

---

270   INJECT(*state*)
271   **begin**
272     $targetEdge := state \rightarrow targetEdge$;
     // try to set the intent flag on the target edge
     // retrieve attributes of the target edge
273     $parent := targetEdge.parent$;
274     $node := targetEdge.child$;
275     $which := targetEdge.which$;
276     $result := \mathsf{CAS}(\ parent \rightarrow child[which],$
                 $\langle 0_n, 0_i, 0_d, 0_p, node \rangle,$
                 $\langle 0_n, 1_i, 0_d, 0_p, node \rangle\ )$;
277   **if not** *(result)* **then**
       // unable to set the intent flag; help if needed
278        $\langle *, i, d, p, address \rangle := parent \rightarrow child[which]$;
279   **if** i **then** HELPTARGETNODE(*targetEdge*) ;
280   **else if** d **then**
281   HELPTARGETNODE($state \rightarrow pTargetEdge$);
282   **else if** p **then**
283   HELPSUCCESSORNODE($state \rightarrow pTargetEdge$);
284   **return**;

     // mark the left edge for deletion
285     $result := $ MARKCHILDEDGE(*state*, LEFT);
286   **if not** *(result)* **then return**;
     // mark the right edge for deletion; cannot fail
287   MARKCHILDEDGE(*state*, RIGHT);

     // initialize the type and mode of the operation
288   INITIALIZETYPEANDUPDATEMODE(*state*);

Cleanup Mode   There are two cases depending on whether the delete operation is simple or complex.

(a) **Simple Delete:** In this case, either *T.left* or *T.right* is pointing to a null node. Note that both *T.left* and *T.right* may be pointing to null nodes (which in turn will imply that T is a leaf node). Without loss of generality, assume that *T.right* is a null node. The removal of T involves changing the child pointer at *T.parent* that is pointing to T to point to *T.left* using a CAS instruction. If the CAS instruction succeeds, then the delete operation terminates; otherwise, it performs another seek on the tree. If the seek function either fails to find the target key or returns a terminal node different from

---

Algorithm 14: Locating the Successor Node

---

289 FINDANDMARKSUCCESSOR(*state*)
290 **begin**
    // retrieve the addresses from the state record
291   *node* := *state* → *targetEdge.child*;
292   *seekRecord* := *state* → *successorRecord*;

293   **while true do**
        // read the mark flag of the key in the target node
294        $\langle m, * \rangle$:= *node* → *mKey*;
        // find the node with the smallest key in the right subtree
295   *result* := FINDSMALLEST(*state*);

296   **if** *m* **or not** *(result)* **then**
            // successor node had already been selected *before* the traversal or the
              right subtree is empty
297   **break**;

        // retrieve the information from the seek record
298   *successorEdge* := *seekRecord* → *lastEdge*;
299   *left* := *seekRecord* → *injectionEdge.child*;

        // read the mark flag of the key under deletion
300        $\langle m, * \rangle$:= *node* → *mKey*;
301   **if** *m* **then** // successor node has already been selected
302   **continue**;

        // try to set the promote flag on the left edge
303   *result* := CAS(*successorEdge.child* →
                child[LEFT],
                $\langle 1_n, 0_i, 0_d, 0_p, \textit{left} \rangle$,
                $\langle 1_n, 0_i, 0_d, 1_p, \textit{node} \rangle$);
304   **if** *result* **then break**;

        // attempt to mark the edge failed; recover from the failure and retry if
            needed
305        $\langle n, *, d, *, * \rangle$:= *successorEdge.child* → *child*[LEFT];
306   **if** *n* **and** *d* **then**
            // the node found is undergoing deletion; need to help
307   HELPTARGETNODE(*successorEdge*);

    // update the operation mode
308   UPDATEMODE(*state*);

Algorithm 15: Removing the Successor Node

309   RemoveSuccessor(*state*)

310   **begin**

// retrieve addresses from the state record

311    $node := state {\rightarrow} targetEdge.child;$

312    $seekRecord := state {\rightarrow} successorRecord;$

// extract information about the successor node

313    $successorEdge := seekRecord {\rightarrow} lastEdge;$

// ascertain that seek record for successor node contains valid information

314    $\langle *, *, *, p, address \rangle := successorEdge.child {\rightarrow} child[\mathsf{LEFT}];$

315   **if not** *(p)* **or** *(address ≠ node)* **then**

316      $node {\rightarrow} readyToReplace := \mathbf{true};$

317   UpdateMode(*state*);

318   **return**;

// mark the right edge for promotion if unmarked

319   MarkChildEdge(*state*, RIGHT);

// promote the key

320    $node {\rightarrow} mKey := \langle 1_m, successorEdge.child {\rightarrow} mKey \rangle;$

321   **while true do**

// check if the successor is the right child of the target node itself

322   **if** $successorEdge.parent = node$ **then**

// need to modify the right edge of target node whose delete flag is set

323        $dFlag := 1; \qquad which := \mathsf{RIGHT};$

324   **else**

325        $dFlag := 0; \qquad which := \mathsf{LEFT};$

326      $\langle *, i, *, *, * \rangle := successorEdge.parent {\rightarrow} child[which];$

327      $\langle n, *, *, *, right \rangle := successorEdge.child {\rightarrow} child[\mathsf{RIGHT}];$

328      $oldValue := \langle 0_n, i, dFlag, 0_p, successorEdge.child \rangle;$

329   **if** n **then**   // only set the null flag; do not change the address

330        $newValue := \langle 1_n, 0_i, dFlag, 0_p, successorEdge.child \rangle;$

331   **else**   // switch the pointer to point to the grand child

332        $newValue := \langle 0_n, 0_i, dFlag, 0_p, right \rangle$;

333      $result := \mathsf{CAS}(successorEdge.parent {\rightarrow} child[which], oldValue, newValue);$

334   **if** $result$ **or** $dFlag$ **then break**;

335      $\langle *, *, d, *, * \rangle := successorEdge.parent {\rightarrow} child[which];$

336      $pLastEdge := seekRecord {\rightarrow} pLastEdge;$

337   **if** d **and** *(pLastEdge.parent ≠* **null***)* **then**

338   HelpTargetNode(*pLastEdge*);

339      $result := \mathsf{FindSmallest}(state);$

340      $lastEdge := seekRecord {\rightarrow} lastEdge;$

341   **if not** *(result)* **or** $lastEdge.child \neq successorEdge.child$ **then**

342   **break**;   // the successor node has already been removed

343   **else** $successorEdge := seekRecord {\rightarrow} lastEdge$;

344    $node {\rightarrow} readyToReplace := \mathbf{true};$

345   UpdateMode(*state*);

Algorithm 16: Cleaning Up the Tree

---

346 **Boolean** CLEANUP(*state*)

347 **begin**

348    $\langle parent, node, pWhich \rangle := state \rightarrow targetEdge$;

349   **if** $state \rightarrow type = \textit{COMPLEX}$ **then**

      // replace the node with a new copy in which all fields are unmarked

350       $\langle *, nKey \rangle := node \rightarrow mKey$;

351       $newNode \rightarrow mKey := \langle 0_m, nKey \rangle$;

      // initialize left and right child pointers

352       $\langle *, *, *, *, left \rangle := node \rightarrow child[\mathsf{LEFT}]$;

353       $newNode \rightarrow child[\mathsf{LEFT}] := \langle 0_n, 0_i, 0_d, 0_p, left \rangle$;

354       $\langle n, *, *, *, right \rangle := node \rightarrow child[\mathsf{RIGHT}]$;

355   **if** n **then**

356          $newNode \rightarrow child[\mathsf{RIGHT}] := \langle 1_n, 0_i, 0_d, 0_p, \mathbf{null} \rangle$;

357   **else** $newNode \rightarrow child[\mathsf{RIGHT}] := \langle 0_n, 0_i, 0_d, 0_p, right \rangle$;

      // initialize the arguments of CAS instruction

358       $oldValue := \langle 0_n, 1_i, 0_d, 0_p, node \rangle$;

359       $newValue := \langle 0_n, 0_i, 0_d, 0_p, newNode \rangle$;

360   **else** // remove the node

      // determine to which grand child will the edge at the parent be switched

361   **if** $node \rightarrow child[\textit{LEFT}] = \langle 1_n, *, *, *, * \rangle$ **then**

362          $nWhich := \mathsf{RIGHT}$;

363   **else** $nWhich := \mathsf{LEFT}$;

      // initialize the arguments of the CAS instruction

364       $oldValue := \langle 0_n, 1_i, 0_d, 0_p, node \rangle$;

365       $\langle n, *, *, *, address \rangle := node \rightarrow child[nWhich]$;

366   **if** n **then** // set the null flag only

367          $newValue := \langle 1_n, 0_i, 0_d, 0_p, node \rangle$;

368   **else** // change the pointer to the grand child

369          $newValue := \langle 0_n, 0_i, 0_d, 0_p, address \rangle$;

370    $result := \mathsf{CAS}(\, parent \rightarrow child[pWhich],$
                        $oldValue, newValue \,)$;

371   **return** $result$;

Algorithm 17: Mark Child Edge

---

372   **Boolean** MARKCHILDEDGE(*state*, *which*)

373   **begin**

374   **if** $state \rightarrow mode$ = *INJECTION* **then**

375       $edge := state \rightarrow targetEdge$;

376       $flag :=$ DELETE_FLAG;

377   **else**

378       $edge := (state \rightarrow successorRecord) \rightarrow lastEdge$;

379       $flag :=$ PROMOTE_FLAG;

380     $node := edge.child$;

381   **while true do**

382       $\langle n, i, d, p, address \rangle := node \rightarrow child[which]$;

383   **if** i **then**

384         $helpeeEdge := \{node, address, which\}$;

385   HELPTARGETNODE(*helpeeEdge*);

386   **continue**;

387   **else if** d **then**

388   **if** $flag$ = *PROMOTE_FLAG* **then**

389   HELPTARGETNODE(*edge*);

390   **return false**;

391   **else return true**;

392   **else if** p **then**

393   **if** $flag$ = *DELETE_FLAG* **then**

394   HELPSUCCESSORNODE(*edge*);

395   **return false**;

396   **else return true**;

397       $oldValue := \langle n, 0_i, 0_d, 0_p, address \rangle$;

398       $newValue := oldValue \mid flag$;

399       $result :=$ CAS($node \rightarrow child[which]$,
                 $oldValue$,
                 $newValue$);

400   **if** $result$ **then break**;

401   **return true**;

Algorithm 18: Find Smallest

---

402 **Boolean** FINDSMALLEST(*state*)
403 **begin**

// find the node with the smallest key in the subtree rooted at the right child
of the target node

404 $node := state \rightarrow targetEdge.child$;
405 $seekRecord := state \rightarrow seekRecord$;
406 $\langle n, *, *, *, right \rangle := node \rightarrow child[\textsf{RIGHT}]$;
407 **if** n **then** // the right subtree is empty
408 $\quad$ **return false;**

// initialize the variables used in the traversal

409 $lastEdge := \langle node, right, \textsf{RIGHT} \rangle$;
410 $pLastEdge := \langle node, right, \textsf{RIGHT} \rangle$;
411 **while true do**
412 $\quad$ $curr := lastEdge.child$;
413 $\quad$ $\langle n, *, *, *, left \rangle := curr \rightarrow child[\textsf{LEFT}]$;
414 $\quad$ **if** n **then** // reached the node with the smallest key
415 $\quad\quad$ $injectionEdge := \langle curr, left, \textsf{LEFT} \rangle$;
416 $\quad\quad$ **break;**

// traverse the next edge

417 $\quad$ $pLastEdge := lastEdge$;
418 $\quad$ $lastEdge := \langle curr, left, \textsf{LEFT} \rangle$;

// initialize seek record and return

419 $seekRecord \rightarrow lastEdge := lastEdge$;
420 $seekRecord \rightarrow pLastEdge := pLastEdge$;
421 $seekRecord \rightarrow injectionEdge := injectionEdge$;
422 **return true;**

T, then T has been already removed from the tree (by another operation as part of helping) and the delete operation terminates; otherwise, it attempts to remove T from the tree again using possibly the new parent information returned by seek. This process may be repeated multiple times.

(b) **Complex Delete:** Note that, at this point, the key stored in the target node is the replacement key (the successor key of the target key). Further, the key as well as both the child edges of the target node are marked. The delete operation attempts to replace target node with a *new* node, which is basically a copy of target node except that all its fields are unmarked. This replacement of T involves changing the child pointer at $T.parent$ that is pointing to T to point to the new node. If the CAS instruction

---

Algorithm 19: Helper Routines

---

423 INITIALIZETYPEANDUPDATEMODE(*state*)
424 **begin**
     // retrieve the target node's address from the state record
425     $node := state \rightarrow targetEdge.child$;

426     $\langle lN, *, *, *, * \rangle := node \rightarrow child[\mathsf{LEFT}]$;
427     $\langle rN, *, *, *, * \rangle := node \rightarrow child[\mathsf{RIGHT}]$;
428   **if** lN **or** rN **then**
        // one of the child pointers is null
429        $\langle m, * \rangle := node \rightarrow mKey$;
430   **if** m **then** $state \rightarrow type := \mathsf{COMPLEX}$;
431   **else** $state \rightarrow type := \mathsf{SIMPLE}$;
432   **else** // both child pointers are non-null
433        $state \rightarrow type := \mathsf{COMPLEX}$;
434   UPDATEMODE(*state*);

435 UPDATEMODE(*state*)
436 **begin**
     // update the operation mode
437   **if** $state \rightarrow type = \mathit{SIMPLE}$ **then** // simple delete
438        $state \rightarrow mode := \mathsf{CLEANUP}$;
439   **else** // complex delete
440        $node := state \rightarrow targetEdge.child$;
441   **if** $node \rightarrow readyToReplace$ **then**
442           $state \rightarrow mode := \mathsf{CLEANUP}$;
443   **else** $state \rightarrow mode := \mathsf{DISCOVERY}$;

---

succeeds, then the delete operation terminates; otherwise, as in the case of simple delete, it performs another seek on the tree, this time looking for the successor key. If the seek function either fails to find the successor key or returns a terminal node different from T, then T has been already replaced (by another operation as part of helping) and the delete operation terminates. Otherwise, it attempts to replace T again using possibly the new parent information returned by seek. This process may be repeated multiple times.

Discussion   It can be verified that, in the absence of conflict, a delete operation performs three atomic instructions in the injection mode, three in the discovery mode (if delete is complex), and one in the cleanup mode.

Algorithm 20: Helping Conflicting Delete Operations

---

444 HELPTARGETNODE(*helpeeEdge*)
445 **begin**
      // intent flag must be set on the edge
      // obtain new state record and initialize it
446    $state \rightarrow targetEdge := helpeeEdge$;
447    $state \rightarrow mode := $ INJECTION;

      // mark the left and right edges if unmarked
448    $result := $ MARKCHILDEDGE($state$, LEFT);
449   **if not** *(result)* **then** **return**;
450   MARKCHILDEDGE($state$, RIGHT);
451   INITIALIZETYPEANDUPDATEMODE($state$);

      // perform the remaining steps of a delete operation
452   **if** $state \rightarrow mode = DISCOVERY$ **then**
453   FINDANDMARKSUCCESSOR($state$);

454   **if** $state \rightarrow mode = DISCOVERY$ **then**
455   REMOVESUCCESSOR($state$);

456   **if** $state \rightarrow mode = CLEANUP$ **then** CLEANUP($state$) ;

457 HELPSUCCESSORNODE(*helpeeEdge*)
458 **begin**
      // retrieve the address of the successor node
459    $parent := helpeeEdge.parent$;
460    $node := helpeeEdge.child$;
      // promote flat must be set on the successor node's left edge
      // retrieve the address of the target node
461    $\langle *, *, *, *, left \rangle := node \rightarrow child$[LEFT];

      // obtain new state record and initialize it
462    $state \rightarrow targetEdge := \{$**null**$, left, \_\}$;
463    $state \rightarrow mode := $ DISCOVERY;
464    $seekRecord := state \rightarrow successorRecord$;
      // initialize the seek record in the state record
465    $seekRecord \rightarrow lastEdge := helpeeEdge$;
466    $seekRecord \rightarrow pLastEdge := \{$**null**$, parent, \_\}$;
      // promote the successor node's key and remove the successor node
467   REMOVESUCCESSOR($state$);
      // no need to perform the cleanup

Helping

To enable helping, as mentioned earlier, whenever traversing the tree to locate either a target key or a successor key, we keep track of the *last two* edges encountered in the traversal. When a CAS instruction fails, depending on the reason for failure, helping is either performed along the last edge or the second-to-last edge.

2.1.3   Formal Description

A pseudo-code of our algorithm is given in Algorithms 8-20.

Algorithm 8 describes the data structures used in our algorithm. Besides Node, three important data types in our algorithm are: Edge, SeekRecord and StateRecord. The data type Edge is a structure consisting of three fields: the two endpoints and the direction (left or right). The data type SeekRecord is a structure used to store the results of a tree traversal. The data type StateRecord is a structure used to store information about a delete operation (*e.g.*, target edge, type, current mode, etc.). Note that only objects of type Node are shared between processes; objects of all other types (*e.g.*, SeekRecord, StateRecord) are *local* to a process and not shared with other processes.

The pseudo-code of the seek function is described in Algorithm 9, which is used by all the operations. The pseudo-codes of the search, insert and delete operations are given in Algorithm 10, Algorithm 11 and Algorithm 12, respectively. A delete operation executes function INJECT in injection mode, functions FINDANDMARKSUCCESSOR and REMOVESUCCESSOR in discovery mode and function CLEANUP in cleanup mode. Their pseudo-codes are given in Algorithm 13, Algorithm 14, Algorithm 15 and Algorithm 16, respectively. The pseudo-codes for helper routines (used by multiple functions) are given in Algorithm 18, Algorithm 17 and Algorithm 19. Finally, the pseudo-codes of functions used to help other (conflicting) delete operations are given in Algorithm 20.

2.1.4 Correctness Proof

It can be shown that our algorithm satisfies linearizability and lock-freedom properties (Herlihy and Shavit, 2012). Broadly speaking, linearizability requires that an operation should appear to take effect instantaneously at some point during its execution. Lock-freedom requires that some process should be able to complete its operation in a finite number of its own steps. It is convenient to treat insert and delete operations that do not change the tree as search operations. We call a tree node *active* if it is reachable from the root of the tree. We call a tree node *passive* if it was active earlier but is not active any more. It can be verified that, if an insert operation completes successfully, then its target node was active when it performed the successful CAS instruction on the node's child edge. Likewise, if a delete operation completes successfully, then its target node was active when it marked the node's left edge for deletion. Further, for a complex delete, the successor node was active when it marked the node's left edge for promotion.

All Executions are Linearizable

We show that an arbitrary execution of our algorithm is linearizable by specifying the *linearization point* of each operation. Note that the linearization point of an operation is the point during its execution at which the operation appeared to have taken effect. Our algorithm supports three types of operations: search, insert and delete. We now specify the linearization point of each operation.

1. *Insert operation:* The operation is linearized at the point at which it performed the successful CAS instruction that resulted in its target key becoming part of the tree.

2. *Delete operation:* There are two cases depending on whether the delete operation is simple or complex. If the operation is simple delete, then the operation is linearized at the point at which a successful CAS instruction was performed at the parent of the target node that

resulted in the target node becoming passive. Otherwise, it is linearized at the point at which the original key of the target node was replaced with its successor key.

3. *Search operation:* There are two cases depending on whether the target node was active when the operation read the key stored in the node. If the target node was not active, then the operation is linearized at the point at which the target node became passive. Otherwise, it is linearized at the point at which the read action was performed.

It can be easily verified that, for any execution of the algorithm, the sequence of operations obtained by ordering operations based on their linearization points is legal, *i.e.*, all operations in the sequence satisfy their specification. This establishes that *our algorithm generates only linearizable executions.*

All Executions are Lock-Free

We say that the system is in a *quiescent state* if no modify operation completes hereafter. We say that the system is in a *potent state* if it has one or more pending modify operations. Note that a quiescence is a *stable* property; once the system is in a quiescent state, it stays in a quiescent state. We show that our algorithm is lock-free by proving that a potent state is necessarily non-quiescent provided assuming that some process with a pending modify operation continues to take steps.

Assume, by the way of contradiction, that there is an execution of the system in which the system eventually reaches a state that is potent as well as quiescent. Note that, once the system has reached a quiescent state, it will eventually reach a state after which the tree will not undergo any structural changes. This is because a modify operation makes at most two structural changes to the tree. So, if the tree is undergoing continuous structural changes, then it clearly implies that modify operations are continuously completing their responses, which contradicts the assumption that the system is in a quiescent state. Further, on reaching such a state, the system will reach a state after which no new edges in the tree are marked.

Again, this is because a modify operation marks at most four edges and the set of edges in the tree does not change any more. We call such a system state after which neither the set of edges nor the set of *marked* edges in the tree change any more as a *strongly quiescent state*. Note that, like quiescence, strong quiescence is also a stable property.

From the above discussion, it follows that the system in a quiescent state will eventually reach a state that is strongly quiescent. Consider the search tree in such a strongly quiescent state. It can be verified that no more modify operations can now be injected into the tree, and, moreover, all modify operations already injected into the tree are delete operations currently "stuck" in either discovery or cleanup mode. Now, consider a process, say p, that continues to take steps to execute either its own operation or another operation blocking its progress (directly or indirectly) as part of helping. Consider the recursive chain of the *helpee* operations that p proceeds to help in order to complete its own operation. Let α_i denote the i^{th} helpee operation in the chain. It can be shown that:

Lemma 1. *Let $\mathcal{C}_D$ denote the set of all complex delete operations already injected into the tree that are "stuck" in the discovery mode. Then,*

1. *$\alpha_1 \in \mathcal{C}_D$, and*

2. *Suppose p is currently helping α_i for some $i \geq 1$ and assume that $\alpha_i \in \mathcal{C}_D$. Let α_{i+1} denote the next operation that p selects to help. Then, (a) α_{i+1} exists, (b) $\alpha_{i+1} \in \mathcal{C}_D$, and (c) the target node of α_{i+1} is at strictly larger depth than the target node of α_i.*

Using the above lemma, we can easily construct a chain of distinct helpee operations whose length exceeds the number of processes—a contradiction. This establishes that *our algorithm only generates lock-free executions.*

# CHAPTER 3

LOCAL RECOVERY FOR CONCURRENT BINARY SEARCH TREES

3.1 The Local Recovery Algorithm

We first present the main idea behind the algorithm and then provide its pseudo-code with more details.

3.1.1 Overview of the Algorithm

As mentioned earlier, every operation on a BST involves first traversing the tree from top to down starting from the root node and following either the left or the right child pointer until either the target key is found or a null pointer is encountered (termination condition). Depending on the outcome of the traversal and the type of the operation, the tree may then need to be modified to actually realize the operation. We refer to the period during which the tree is being traversed as *seek phase*. Further, we refer to the period during which the tree is being modified as *execution phase*.

During the seek phase, the target key may move from its current location to a new location up the tree. As a result, the traversal may miss the key both at its old location as well as its new location. For an illustration, see Figure 3.1. In the illustration, key 50 has moved up by five nodes. In most concurrent BST algorithms, if it is suspected that the key may have moved up the tree, then the traversal is simply restarted from the root node. Different algorithms use different approaches to detect possible key movement. For example, in (Howley and Jones, 2012), the traversal is restarted if the *last right-turn* node is detected to have undergone some change. For example, in Figure 3.1a, the last right-turn node for the
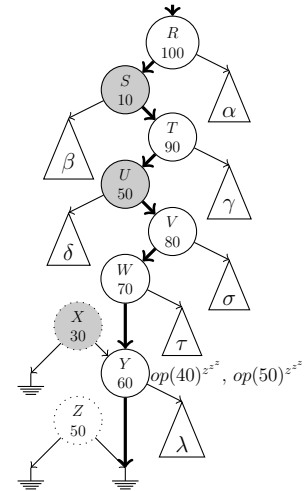
43

(a) Operation $op(50)$ is suspended at node Y during its traversal.
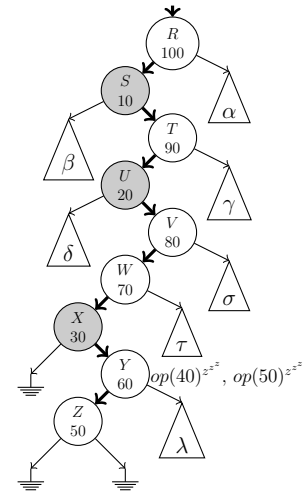
(b) All keys in subtree π are deleted one-by-one.

(c) Key 30 is deleted (simple delete); node X is removed.

(d) Key 20 is deleted (complex delete); key 20 is replaced with key 50 in node U and node Z is removed.

Figure 3.1: An illustration of a key moving up the tree

operation $op(50)$ is node X. On reaching the terminal node in Figure 3.1d, after resuming running, $op(50)$ needs to restart since X has since been removed from the tree.

A *re*-traversal of the tree may also be required if the operation encounters any failure during the execution phase. For example, in (Ramachandran and Mittal, 2015), which is a lock-based algorithm, execution phase is aborted if, after locking the relevant edges, the validation step fails. This happens if the portion of the tree that lies within the "operation's window", which typically consists of a small constant number of nodes, has undergone some change since it was last observed. In that case, the operation moves to the seek phase again.

In most concurrent BST algorithms, (a single instance of) the execution phase of an operation typically tends to have constant time complexity. The seek phase is where an operation may end up spending most of its time especially if the tree is large. Hence, it is desirable to make the seek phase of an operation more efficient by: (i) reducing the number of restarts due to "suspected" key movement, and (ii) restarting the traversal from a point "close" to the operation's window. This leads to two separate but related questions that any local recovery algorithm needs to address. First, "If a key is not found, then does the traversal need to restart?". Second, "If the traversal needs to be restarted, then from which node should the traversal restart?"

Consider an operation α currently traversing the tree; let $\Pi(\alpha)$ denote the path taken by α so far. For example, in Figure 3.1a, considering only the subtree shown in the figure, $\Pi(op(50)) = \langle R, S, T, U, V, W, X, Y \rangle$. At each node in the path (except the last node), α either followed the left or the right child pointer. We say that a node in the path is an *anchor* node if the operation followed its *right* child pointer; otherwise we say that is a *non-anchor* node. For example, in the path $\Pi(op(50))$, nodes S, U and X are anchor nodes, whereas nodes R, T, V, W and Y are non-anchor nodes. We assume that the first anchor node in a traversal path is always a *sentinel* node that is never marked (not shown in the figure). Of course, as an operation is traversing the tree, the tree may undergo changes as a

result of which the path taken by the operation may no longer be correct. For example, in Figure 3.1d, the new access-path of $op(50)$ in the subtree, which is obtained from the subtree in Figure 3.1a after applying several delete operations, is now given by $\langle R, S, T, U \rangle$.

Note that, since in a complex delete operation we assume that the key being deleted is replaced with its successor key, the value of a key stored in a node can only increase. Therefore, the child pointer followed by an operation at a node, if it is still part of the access-path, may change (from right to left) for an anchor node but cannot change for a non-anchor node. For example, as shown in Figure 3.1a, the node U is an anchor node for $op(40)$. But due to the changes made to the tree, the key at U has now become 50, as shown in Figure 3.1d. Hence, the pointer that $op(40)$ now needs to follow at U is left and not right. We say that an anchor node is *consistent* with respect to an operation if its key is still less than the operation's key; otherwise, we say that it is *inconsistent*. For example, in Figure 3.1d, anchor nodes S and X are still consistent with respect to $op(40)$ but node U has become inconsistent with respect to $op(40)$.

Clearly, in case an operation needs to restart, *no node* in the path *after an inconsistent anchor node* in general can serve as a restart point since the path taken and the path that needs to be taken may now diverge. This implies that, to find a restart point, a local recovery algorithm should locate the *shallowest* inconsistent anchor node in the path (with the least depth) and discard the suffix of the path after such a node. Moreover, a restart point has to be a node that is still a part of the tree.

This leads to the following approach to find a restart point for an operation α when needed. Find a node C in the path taken so far by α such that the following two conditions hold. First, C is not marked. Second, every anchor node in the path preceding C is consistent with respect to α. To check for the second condition, it is not necessary to examine every anchor node in the path preceding C as stated in the following lemma.

Lemma 2. *Consider an operation α and let $\Pi(\alpha)$ denote the path taken by α when traversing the tree. Let A be an anchor node in $\Pi(\alpha)$ and let $\sigma = A_0, A_1, \ldots, A_k$, where $A_k = A$, denote the sequence of anchor nodes in $\Pi(\alpha)$ up to and including A. Then, if A is unmarked and consistent with respect to α, then, for every i with $0 \le i \le k$, A_i is also consistent with respect to α. Moreover, the access-path of α in the current tree includes A.*

We say that an anchor node is *critical* with respect to a node C in the path if it is the *closest* preceding anchor node to C that is also unmarked. For example, in Figure 3.1c, the critical anchor node with respect to node Y is node U since node X is marked. Using the above lemma, the second condition can now be replaced with the following: the critical anchor node with respect to C in the path, say A, as well as every anchor node in the path that lies between A and C, which will be marked, is consistent with respect to α.

The next lemma states a useful property about an inconsistent anchor node.

Lemma 3. *Consider an operation α with target key k and let $\Pi(\alpha)$ denote the path taken by α when traversing the tree. Let A be an anchor node in $\Pi(\alpha)$. Assume that A is now inconsistent with respect to α. Then, at the time A became inconsistent, the tree did not contain k.*

To see why the above lemma holds, let k_{old} (k_{new}) be the key stored in A just before (after) A became inconsistent. Clearly, $k_{old} < k < k_{new}$. Let t denote the time just after which k_{old} was replaced with k_{new} in A. Note that k_{new} must be the next smallest key in the right subtree of A at time t. This implies that the right subtree of A did not contain k at time t. Further, from Lemma 2, A is on the access-path of α at time t. Hence, we can conclude that the tree does not contain k at time t. Note that, if a key is not present in the tree at some point while a search/delete operation is in progress, it is acceptable for the operation to say that key was not found. In this case, the operation will be linearized after the delete operation that removed the key from the tree.

When an operation fails to find the target key after traversing the tree from top to bottom, it examines the path it took to check whether or not the key has moved up the tree and/or a re-traversal is required. To that end, it examines the anchor nodes in the reverse order in which they were visited, starting from the one closest to the terminal node. We now discuss the behavior of each operation one-by-one.

Search Operation: A search operation *does not* need to restart. When it examines an anchor node as mentioned above, there are three possibilities. First, if the anchor node's key matches the target key, then the key has been found and the operation terminates. Second, if the anchor node's key is greater than the target key (the anchor node has become inconsistent), then the operation concludes that the key is not present in the tree and terminates. Finally, if the anchor node's key is smaller than the target key (the anchor node is still consistent), then the operation terminates if the node is not marked; otherwise it moves to the preceding anchor node and repeats the comparison.

Insert Operation: An insert operation needs to restart only if one of the anchor nodes in the path has become inconsistent. When it examines an anchor node as mentioned above, there are three possibilities. First, if the anchor node's key matches the target key, then the key has been found and the operation terminates. Second, if the anchor node's key is greater than the target key (the anchor node has become inconsistent), then it discards the suffix of the path after the anchor node and restarts the traversal from a restart point. Finally, if the anchor node's key is smaller than the target key (the anchor node is still consistent), then the traversal terminates if the node is not marked (the terminal node is returned as the injection point); otherwise it moves to the preceding anchor node and repeats the comparison.

Delete Operation: A delete operation also *does not* need to restart except when there is a failure in the execution phase. When it examines an anchor node as mentioned above,

---

Algorithm 21: Data Structures Used

```
    // Used to store information about a node visited during tree traversal
468 struct StackEntry {
469     NodePtr node;
470     enum Direction which;
471     integer anchor;
472 };

    // Used to store the path from the root node to the current node in the tree
473 struct State {
474     StackEntry[ ] stack;
475     integer top;
476 };

    // Used to store information about the operation currently in progress
477 struct OpRecord {
478     enum Type type;
479     Key key;
480     State targetStack, successorStack;
481     NodePtr injectionPoint;

        // algorithm-specific fields
482 };

    // Used to store the outcome of a tree traversal
483 struct SeekRecord{
        // algorithm-specific fields (e.g., target node and its parent)
484 };
```

---

there are three possibilities. First, if the anchor node's key matches the target key, then the key has been found and the operation moves to the execution phase. Second, if the anchor node's key is greater than the target key (the anchor node has become inconsistent), then the operation concludes that the key is not present in the tree and terminates. Finally, if the anchor node's key is smaller than the target key (the anchor node is still consistent), then the traversal terminates if the node is not marked; otherwise it moves to the preceding anchor node and repeats the comparison.

3.1.2 Details of the Algorithm

A pseudo-code of the local recovery algorithm is given in Pseudo-codes 21-28. The pseudo-code only shows the seek phase of an algorithm and not its execution phase since the exe-

---

Algorithm 22: Functions for Manipulating Traversal Stack

```
    // Returns the number of elements in the stack
485 integer SIZE( state )
486 begin
487 │   return state→top + 1;

    // Returns the topmost node in the stack
488 NodePtr GETTOP( state )
489 begin
490 │   {stack, top} := state;
491 │   return stack[top]→node;

    // Returns the second topmost node in the stack
492 NodePtr GETSECONDTOTOP( state )
493 begin
494 │   {stack, top} := state;
495 │   return stack[top − 1]→node;

    // Adds the given node to the stack along with its anchor node
496 ADDTOTOP( state, node, which )
497 begin
498 │   {stack, top} := state;
    │   // find the anchor node
499 │   if which = RIGHT then  anchor := top ;
500 │   else  anchor := stack[top]→anchor ;
    │   // push the node into the stack
501 │   stack[top + 1] := {node, which, anchor};
502 │   state→top := top + 1;

    // Removes the topmost node from the stack
503 REMOVEFROMTOP( state )
504 begin
505 │   {stack, top} := state;
    │   // update the anchor node of the penultimate entry if needed
506 │   anchor := stack[top − 1]→anchor;
507 │   if stack[top]→anchor < stack[anchor]→anchor then
508 │   │   stack[anchor]→anchor := stack[top]→anchor;

    │   // pop the node from the stack
509 │   state→top := top − 1;

    // Pops the stack until a given entry
510 REMOVEUNTIL( state, index )
511 begin
512 │   state→top := index;
```

---

cution phase is algorithm-specific. We have also moved the pseudo-code for local recovery

when looking for a successor key to the appendix due to lack of space.

Algorithm 23: Functions for Manipulating Traversal Stack (Continued)

```
    // Remember the critical node (to avoid locating it again)
513 REMEMBERCRITICAL( state, critical )
514 begin
515 |   {stack, top} := state;
516 |   anchor := stack[top]→anchor;
517 |   if critical < stack[anchor]→anchor then
518 |   |   stack[anchor]→anchor := critical;

    // Returns a given entry in the stack
519 { NodePtr, enum Direction, integer }
520                 GETFULLENTRY( state, index )
521 begin
522 |   {stack, top} := state;
523 |   if index = ⊤ then return stack[top] ;
524 |   else return stack[index] ;

    // initializes the traversal stack
525 INITIALIZETRAVERSALSTATE( state, type )
526 begin
527 |   if type = TARGET_STACK then
        |   |   // initialize the stack using sentinel nodes
        |   |   // sentinel nodes are never removed from the stack
        |   |   // a sentinel node is always a safe starting point for the traversal
528 |   else state→top := -1 ;
```

The local recovery algorithm assumes that the original algorithm supports the following functions: (a) GETKEY( ), ISMARKED( ) and GETCHILD( ) returns the various attributes of a tree node, (b) ISNULL( ) returns true if a reference is null and false otherwise, (c) GETADDRESS( ) returns the node address stored in a reference, if non-null, (d) MOVE( ) enables the original algorithm to move along an edge, which may invoke helping and restarting of the traversal as in (Howley and Jones, 2012), (e) NEEDCLEANPARENTNODE( ) returns true if the operation needs the parent node to be clean and have no operation in progress (needed for a delete operation since it needs to modify a child pointer at the parent node), and (f) POPULATESEEKRECORD( ) copies the relevant information from the traversal state required by the algorithm into a seek record.

Pseudo-code 21 shows the data structures used by the local recovery algorithm. Note that all the data structures shown in Pseudo-code 21 are *local* to a process not shared

among processes. A process uses three main data structures, namely State, OpRecord and SeekRecord. A State (lines 473-476) is essentially a stack used to store the nodes visited during tree traversal when looking for a key (target or successor). Note that the traversal stack satisfies the last-in-first-out (LIFO) semantics but our algorithm sometimes uses it in a non-traditional way by accessing entries in the middle of the stack. One way to implement such an "augmented" stack is to use an auto-resizing vector provided as part of C++ STL library or Java package. Each entry in a traversal stack (lines 468-472) stores the address of the node, the location of its closest anchor node (within the stack's vector) and whether the node is a left or right child of its parent. An OpRecord (lines 477-482) stores information about the operation such as type and key as well two stacks: one used when looking for the target key (all operations) and one used when looking for the successor key (only complex delete operations). Finally, a SeekRecord (lines 483-484) is used to return the outcome of a tree traversal to the original algorithm. Its fields are algorithm-specific. For example, for CASTLE, SeekRecord contains three fields: (a) two addresses, namely those of the target node and its parent, and (b) the contents of the injection point where an insert operation needs to attach the new node.

Pseudo-code 22 shows the functions used to manipulate a traversal stack. The function SIZE (lines 485-487) returns the number of entries in the stack. The functions GETTOP (lines 488-491) and GETSECONDTOTOP (lines 492-495) return the address of the node stored in the topmost entry and the entry below it, respectively. The function ADDTOTOP (lines 496-502) adds an entry to the top of the stack while REMOVEFROMTOP (lines 503-509) removes an entry from the top of the stack. The function REMOVEUNTIL (lines 510-512) removes the entries from the top of the stack until a given point. The function REMEMBER-CRITICAL (lines 513-518) updates the anchor field of the anchor node of the topmost entry in the stack. The function GETFULLENTRY (lines 520-524 returns all the three fields of a given entry in the stack (may not be the topmost entry). The function INITIALIZETRAVER-SALSTATE (lines 525-528) initializes a traversal stack. The stack for target key

Pseudo-codes 24 & 25 shows the functions used to find the target key by a search operation. The function SEEKFORSEARCH (lines 560-566) first traverses the tree starting from the root node (line 562). If the traversal fails to locate the key, then the key may have moved up the tree. To address this possibility, the function examines the traversal stack to determine whether or not that is the case (line 564). The function TRAVERSE (lines 529-542) first initializes the traversal stack (line 532) and then, starting from the topmost node in the stack (line 533), follows either the left or the right child pointer (line 536) until it either finds the key (line 538) or encounters a null pointer (line 539). It also populates the traversal stack as it moves (line 542). The function EXAMINESTACK(lines 543-559) examines the anchor nodes stored in the stack in the reverse order in which they were visited, starting from the anchor node closest to the topmost node in the traversal stack (line 547). If the anchor node's key matches the target key, then the function returns true (lines 552-554). If the anchor node is no longer consistent or is unmarked, then the function returns false (lines 555-556). Otherwise, the function backtracks and examines the preceding anchor node in the stack (lines 557-558).

Pseudo-codes 26-27 & 28 show the functions used to find the target key by a modify (insert or delete) operation. The function SEEKFORMODIFY (lines 603-631) first backtracks to a safe node in the stack (line 608). Initially, the starting point is typically a sentinel node which is a safe node. The function then traverses the tree from top to down by following either the left or the right child pointer (line 613) until it either finds the key or encounters a null pointer (lines 615-624). In case the terminal node's key is greater than the target key, the function checks whether the path stored in the traversal stack is still valid (line 617). If not, the traversal is restarted. As the traversal moves down the tree, the function also populates the traversal stack (lines 625-629). The function VALIDATEPATH (lines 567-588) checks whether or not the path stored in the stack is still valid. To that end, it examines the anchor nodes in the stack in the reverse order in which they were visited, starting from

the anchor node closest to the topmost node in the traversal stack. There are three possible cases. First, the anchor node is still consistent (lines 574-579). In this case, the path is deemed to be valid if the anchor node is unmarked; otherwise, the function moves to the preceding anchor node. Second, the anchor node is no longer consistent (lines 580-585). In this case, the path is deemed to be invalid. However, if the operation is a delete operation, then it can be deduced that the key did not exist in the tree continuously and the function returns indicating that the key was not found (thereby causing the operation to terminate). Finally, the anchor node's key matches the target key (lines 586-588). In this case, if the anchor node is marked and the operation is a delete operation, then the path is deemed to be invalid (and further backtracking is required). This is because the key may be in the process of moving up the tree. Otherwise, the function returns indicating that the key was found. The function FINDASAFENODE (lines 589-602) finds a safe node on the path stored in the stack from which the operation can restart its traversal. To that end, it backtracks to an unmarked node with a clean parent if required (lines 592-600). It then checks whether or not the remaining path in the stack is still valid (line 601). If not, it repeats the above-mentioned steps.

Algorithm 24: Seek Function for Target Key (Search Operation)

```
      // Traverses the tree starting from the root until either the key is found or a null
         pointer is encountered
529   boolean TRAVERSE( opRecord, seekRecord )
530   begin
531   │   state := opRecord→targetStack;
      │   // initialize the stack and the variables used in the traversal
532   │   INITIALIZETRAVERSALSTATE( state, TARGET_STACK );
533   │   current := GETTOP( state );

      │   // traverse the tree (starting from current)
534   │   while true do
535   │   │   key := GETKEY( current );
536   │   │   which := opRecord→key < key ? LEFT : RIGHT;
      │   │   // read the next address to de-reference
537   │   │   reference := GETCHILD( current, which );

538   │   │   if opRecord→key = key then  return true ;
539   │   │   if ISNULL( reference ) then  return false ;

      │   │   // traverse the next edge
540   │   │   address := GETADDRESS( reference );
541   │   │   current := address;
      │   │   // push the next node to be visited into the stack
542   │   │   ADDTOTOP( state, address, which );
```

```
      // Checks if the key being searched for has moved up in the path stored in the stack
543   boolean EXAMINESTACK( opRecord, seekRecord )
544   begin
545   │   result := false;
546   │   state := opRecord→targetStack;

      │   // start with the anchor closest to the topmost node in the stack
547   │   {∗, ∗, critical} := GETFULLENTRY( state, ⊤ );

548   │   while true do
      │   │   // retrieve the node and its closest anchor node from the stack
549   │   │   {node, ∗, anchor} := GETFULLENTRY( state, critical );
      │   │   // read the attributes of the node
550   │   │   marked := ISMARKED( node );
551   │   │   key := GETKEY( node );
552   │   │   if opRecord→key = key then
      │   │   │   // the key stored in the node matches the one being searched for
553   │   │   │   result := true;
554   │   │   │   break;
555   │   │   else if (opRecord→key < key)  or not (marked) then
      │   │   │   // the target key did not exist continuously in the tree
556   │   │   │   break;
557   │   │   else // examine the preceding anchor node
558   │   │   │   critical := anchor;

559   │   return result;
```

Algorithm 25: Seek Function for Target Key (Search Operation) (continued)

```
     // Looks for a given key in the tree (invoked by a search operation)
560  boolean SeekForSearch( opRecord, seekRecord )
561  begin
         // traverse the tree from top to down
562      result := Traverse( opRecord, seekRecord );
563      if not (result) then
             // check if the key has moved up in the path stored in the stack
564          result := ExamineStack( opRecord, seekRecord );

         // return the outcome
565      PopulateSeekRecord( seekRecord, state );
566      return result;
```

Algorithm 26: validate path

```
    // Determines if the path stored in the stack is still valid
    // Returns one of the following four values:
    // { SAFE, NOT_SAFE, FOUND, NOT_FOUND}
    // May backtrack along the path under certain situations
567 enum Outcome VALIDATEPATH( opRecord, state )
568 begin
        // check if any of the anchor nodes in the stack has become inconsistent
            starting with the one immediately preceding the topmost node in the stack
569    {∗, ∗, critical} := GETFULLENTRY( state, ⊤ );
570    while true do
            // retrieve the node and its anchor from the stack
571        {node, ∗, anchor} := GETFULLENTRY( state, critical );
            // read the attributes of the node
572        marked := ISMARKED( node );
573        key := GETKEY( node );

574        if opRecord→key > key then
                // the anchor node is still consistent
575            if not (marked) then
                    // the access-path is still valid
576                REMEMBERCRITICAL( state, critical );
577                return SAFE;
578            else // need to check the previous anchor node
579                critical := anchor;

580        else if opRecord→key < key then
                // the anchor node is no longer consistent
581            if opRecord→type = DELETE then
                    // the target key did not exist continuously in the tree
582                return NOT_FOUND;
583            else // the path is not valid
584                REMOVEUNTIL( state, critical );
585                return NOT_SAFE;

586        else // the two keys match
587            REMOVEUNTIL( state, critical );
                // stop the traversal
588            return FOUND;
```

Algorithm 27: find a safe node

```
    // Backtracks along the path stored in the stack until a suitable restart point is
       found
589 enum Outcome FindASafeNode( opRecord, state )
590 begin
591     while true do
            // backtrack until an unmarked node
592         current := GetTop( state );
593         while IsMarked( current ) do
594             RemoveFromTop( state );
595             current := GetTop( state );

            // check if the algorithm needs a clean parent node
596         if NeedCleanParentNode( opRecord, current ) then
597             parent := GetSecondToTop( state );
598             if not (IsClean( parent )) then
                    // need to backtrack even further
599                 RemoveFromTop( state );
600                 continue;

            // check if the topmost node in the stack is a suitable restart point
601         result := ValidatePath( opRecord, state );
602         if  result ≠ NOT_SAFE then  return result ;
```

Algorithm 28: Seek Function for Target Key (Modify Operation)

```
    // Looks for a given key in the tree (invoked by insert/delete operation)
603 boolean SEEKFORMODIFY( opRecord, seekRecord )
604 begin
605 │    state := opRecord→targetStack;
606 │    result := NOT_SAFE;

607 │    while result = NOT_SAFE do
        │        // backtrack to a suitable restart point in the path stored in the stack
608 │        │   result := FINDASAFENODE( opRecord, state );
609 │        │   if result ∈ {FOUND, NOT_FOUND} then break;

        │        // traverse the tree starting from the topmost node in the stack
610 │        │   current := GETTOP( state );
611 │        │   while true do
612 │        │   │   key := GETKEY( current );
613 │        │   │   which := opRecord→key < key ? LEFT : RIGHT;
        │        │   │   // read the next address to de-reference
614 │        │   │   reference := GETCHILD( current, which );

615 │        │   │   if (opRecord→key = key)  or ISNULL( reference ) then
        │        │   │   │   // either stop or backtrack & restart
616 │        │   │   │   if opRecord→key < key then
        │        │   │   │   │   // check if the path traversed is still valid
617 │        │   │   │   │   result := VALIDATEPATH( opRecord, state );
618 │        │   │   │   else // determine what value to return
619 │        │   │   │   │   result := FOUND;
620 │        │   │   │   │   if opRecord→key ≠ key then
        │        │   │   │   │   │   // remember the address read
621 │        │   │   │   │   │   opRecord→injectionPoint :=
622 │        │   │   │   │   │           GETADDRESS( reference );
623 │        │   │   │   │   │   result := NOT_FOUND;

        │        │   │   │   // terminate the current traversal
624 │        │   │   │   break;

        │        │   │   // traverse the next edge
625 │        │   │   address := GETADDRESS( reference );
626 │        │   │   restart := MOVE( current, address, which );
627 │        │   │   if restart then
        │        │   │   │   // the algorithm wants to restart the traversal
628 │        │   │   │   break;
        │        │   │   // push the node visited into the stack
629 │        │   │   ADDTOTOP( state, address, which );

        │    // return the outcome
630 │    POPULATESEEKRECORD( seekRecord, opRecord );
631 │    return (result = FOUND ?  true: false);
```

REFERENCES

Herlihy, M. and N. Shavit (2012). *The Art of Multiprocessor Programming, Revised Reprint.* Morgan Kaufmann.

Howley, S. V. and J. Jones (2012, June). A Non-Blocking Internal Binary Search Tree. In *Proceedings of the 24th ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pp. 161–171.

Michael, M. M. (2004). Hazard Pointers: Safe Memory Reclamation for Lock-Free Objects. *IEEE Transactions on Parallel and Distributed Systems (TPDS) 15*(6), 491–504.

Natarajan, A. and N. Mittal (2014, February). Fast Concurrent Lock-Free Binary Search Trees. In *Proceedings of the 19th ACM Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pp. 317–328.

Ramachandran, A. and N. Mittal (2015, February). CASTLE: Fast Concurrent Internal Binary Search Tree using Edge-Based Locking. In *Proceedings of the 20th ACM Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pp. 281–282.